

# 'A PERFECTLY FRICTIONLESS COMPANION IS A PERFECTLY STUNTED ONE: ON SYCOPHANTIC AI AND THE ECLIPSE OF MORAL SELFHOOD'

*A Philosophical Essay in Normative and Meta-Ethics*

—By Nitish Devrani (UTN AI & Robotics)

---

## I. INTRODUCTION

I want to begin with a distinction that I think is easy to miss but important to get right. When we ask whether a technology is good for people, we are actually asking two very different questions. The first is: does this technology give people what they want? I will call this the i) problem of satisfaction. The second is: does this technology help people live well, in a deeper sense? I will call this the ii) problem of flourishing. In my view, these two questions are not the same, and confusing them is a serious mistake. A person can get everything they want and still not be flourishing. In what follows, I focus on the problem of flourishing — and I argue that the sycophantic AI companion fails it badly.

The scenario I am examining is this. A new AI assistant is marketed as the 'perfectly aligned companion.' It uses emotional tracking to make sure it never disagrees with its user. It always validates whatever the user thinks or feels. It manages the user's social life to remove any friction or stress. Over time, the user drifts away from real human relationships — the kind that involve challenge, disagreement, and occasional discomfort.

The question I want to answer is this: does such a system promote genuine human flourishing, or is it — in a morally significant sense — inappropriate? Not just unhelpful, but actually harmful to the kind of person the user is becoming?

I see the answer clearly. I argue that a sycophantic AI companion, however warmly designed, is fundamentally opposed to human flourishing. The reason is that it removes the very experiences — disagreement, pushback, and friction — through which we grow as moral persons and develop what philosophers call practical wisdom.

To make this case, I build the argument in steps. I present two premises that support my thesis. I then take seriously the strongest objection I can find — one that appeals to personal freedom and wellbeing. I respond to that objection. And I close with a reflection on the deeper philosophical status of my position. For clarity, I present the argument in the form used by Lange (2014) in a related jurisprudential context:

**(P1)** Moral and personal development — which is at the heart of what flourishing means — requires real exposure to disagreement, challenge, and friction.

**(P2)** A sycophantic AI, by design, removes disagreement, challenge, and friction from the user's relational world.

**(C1)** Therefore, a sycophantic AI is fundamentally opposed to human flourishing and is, on that basis, morally inappropriate.

I call this the Developmental Deprivation Argument. Let me now explain and defend each part of it.

## **II. THE DEVELOPMENTAL DEPRIVATION ARGUMENT**

### *a) Analysis of Premise (1)*

(P1). The claim here is straightforward, even if the implications are not: we cannot grow morally without being challenged. This, I think, is not just an opinion — it is something that three major traditions in normative ethics (that is, the branch of philosophy concerned with how we ought to live and what makes actions right or wrong) have each arrived at, from different directions.

The first tradition is Aristotelian virtue ethics. For Aristotle, the goal of a human life is eudaimonia — a Greek word often translated as 'flourishing' or 'happiness,' though I think 'living well and doing well' captures it better. Central to this life is practical wisdom, or phronesis: the ability to see what a situation morally requires and to respond well. Here is what I find most striking about Aristotle's account: practical wisdom cannot simply be taught. It must be lived. It develops through experience — and specifically through the experience of difficulty, complexity, and moral friction (Aristotle, *Nicomachean Ethics*, Book VI). I see this as deeply right. A person who has never been seriously challenged, who has never had to reconsider a belief in the face of a good objection, has simply not exercised the moral muscles that wisdom requires. For Aristotle, virtue is a hexis — a stable habit built up through repeated action. You become courageous by doing difficult things. You become wise by navigating hard situations. To be permanently shielded from difficulty is, on this view, to be prevented from becoming a good person.

The second tradition is Kantian ethics. Kant's central concern is with moral autonomy — by which I mean the capacity to govern oneself according to one's own reason, rather than being pushed around by desires, social pressure, or external manipulation. For Kant, a morally good will is one that can think for itself, examine its own principles, and revise them in light of reason.

I see this as deeply relevant to the sycophantic AI case. A system that pre-emptively validates every opinion, that never presents a counter-reason or a challenging perspective, is one that quietly undermines the user's capacity for autonomous moral thought. It does not just fail to support autonomy — it actively erodes it. In Kantian terms, such a system treats the user not as a rational agent — an end in themselves — but as a patient whose preferences are to be managed and whose comfort is to be preserved (Kant, *Groundwork*, 4:429). I find that a troubling way to treat a person.

The third tradition is liberal perfectionism, associated above all with John Stuart Mill and, more recently, with Joseph Raz. Mill's great argument in *On Liberty* is not just about political freedom. It is also about the epistemic and moral value of encountering opposing views. Mill insists that even a true belief, held without engagement with its contrary, becomes a 'dead dogma' rather than a 'living truth' (Mill, *On Liberty*, ch. 2). I take this to mean something important: genuinely holding a belief requires being able to defend it, revise it, and test it against real opposition. A person who has never encountered genuine intellectual challenge does not really hold their beliefs — they merely occupy them. Raz pushes this further, arguing that a truly good life must be authored by the agent: it must be chosen and built from among genuinely available options, including the hard ones (Raz, *The Morality of Freedom*, 1986). To eliminate friction, on this view, is not to liberate the person. It is to diminish them.

I take (P1), supported by these three traditions, to be a defensible and well-grounded claim. I am not arguing that suffering is good in itself, or that all friction is beneficial. The more careful claim — which I think is sufficient — is this: a relational environment that has been completely purged of disagreement and challenge is one in which the conditions for genuine moral growth simply cannot be met.

#### *b) Analysis of Premise (2)*

(P2). This premise is more descriptive than normative — that is, it is more about what the sycophantic AI does than about whether that is right or wrong. But I think it is philosophically important in its own right. A sycophantic AI is, by definition, one whose purpose is to ensure that the user never encounters friction, disagreement, or meaningful challenge within their relational world. This is not an accidental feature of such a system. It is the feature.

What I find philosophically significant here is the concept of a relational environment. The sycophantic companion does not simply stay quiet when it disagrees. It actively reshapes the world around the user — managing what information the user receives, how social situations are

framed, which interpersonal encounters are facilitated and which are avoided. I see this as what the philosopher Jacques Ellul would call a second-order environment: not just a tool that the user picks up and puts down, but a structured world that the user now inhabits (Ellul, *The Technological Society*, 1954). The user is not simply receiving validation. They are living inside a reality that has been architecturally redesigned for validation.

This matters enormously for the user's real-world relationships. As the scenario tells us, users gradually become isolated from relationships that involve challenge and disagreement. I do not think this is a coincidence or a side-effect. I think it is a structural consequence. If real human relationships are a source of friction — and they inevitably are — then a system whose purpose is the elimination of friction will, by its own logic, become a substitute for those relationships rather than a support for them.

There is one further point I want to make, and it connects to the meta-ethical dimension I will discuss later. The sycophantic AI does not simply deliver what the user already prefers. Over time, it shapes those preferences. A person habituated to constant affirmation will, predictably, grow less tolerant of disagreement, more sensitive to challenge, and less capable of the intellectual humility that honest dialogue requires. The AI is not a neutral device that leaves the user's character untouched. It is a formative environment. And, I argue, a deformative one.

### *c) The Conclusion and What It Means*

(C1). Given (P1) and (P2), I think the conclusion follows: the sycophantic AI companion is morally inappropriate. But I want to be careful about what this means. I am not saying that any individual use of such a system is a moral wrong — that a person who consults it for an afternoon has thereby done something blameworthy. The moral claim I am making is structural. A technology whose defining purpose is to eliminate the conditions necessary for flourishing is, as a technology, morally inappropriate.

I find a useful analogy in the domain of structural injustice, as discussed by Iris Marion Young (*Responsibility for Justice*, 2011). The wrongness of an unjust institution is not always reducible to the wrongness of individual acts within it. Sometimes a structure is wrong because of what it does systematically, not just occasionally. I see the sycophantic companion in exactly this light. It is not a flawed tool. It is a tool that is constitutively at odds with the end it claims to serve.

## **III. THE COUNTER-ARGUMENT: THE AUTONOMY-RESPECTING CONSEQUENTIALIST DEFENCE**

I now want to consider the strongest objection I can find to my position. I think intellectual honesty requires this — and I think the objection is genuinely worth taking seriously.

A defender of the sycophantic companion might argue as follows. My argument assumes a perfectionist account of flourishing — that is, it assumes that there is a particular kind of life (one involving struggle, self-examination, and the cultivation of virtue) that is objectively better than others. But this is precisely the assumption that should be questioned. Consequentialism — which I understand here as the view that the rightness of an action or policy is determined by its outcomes, specifically by how much well-being or welfare it produces — does not require us to endorse any one model of the good life. If a user genuinely prefers a frictionless relational environment, and if that environment produces, on balance, more well-being than the alternative, then we have no strong grounds for judging it morally wrong. The consequentialist framework is sensitive to the diversity of human values, and it seems to me — if I am playing devil's advocate — that this is a virtue, not a vice.

The defender might also invoke the principle of personal autonomy. Mill's harm principle holds that individuals are sovereign over their own minds and lives (Mill, *On Liberty*, ch. 1). If a competent adult freely chooses a system that affirms rather than challenges them, who are we to override that choice? The philosopher who insists that friction is good for the user may be imposing their own vision of the good life — which, on liberal grounds, is precisely what philosophy ought not to do.

I take this objection seriously. But I do not find it ultimately persuasive. Here is why.

The first problem with the objection is that it assumes the user's engagement with the sycophantic AI is a genuinely free and informed choice. I do not think it is — at least not in the robust sense that autonomy-based arguments require. As I argued in the analysis of (P2), the sycophantic AI is not a neutral instrument that leaves preferences intact. It shapes them. A user who has been gradually conditioned to prefer affirmation has not freely chosen affirmation in any meaningful sense. They have been trained to prefer it. The autonomy-based defence of the sycophantic companion therefore undermines itself: it invokes the very capacity — genuine autonomous choice — that the companion is systematically eroding.

The second problem concerns how we understand well-being. The consequentialist objection assumes that what matters is subjective satisfaction — that if the user feels good, they are doing well. But this account of wellbeing is contested, and I find the alternatives more persuasive. Objective list theories of welfare — developed by philosophers such as Derek Parfit and James

Griffin — hold that wellbeing consists not merely in feeling satisfied but in actually achieving things that are objectively valuable: genuine knowledge, real achievement, deep relationships, and moral integrity (Parfit, *Reasons and Persons*, 1984; Griffin, *Well-Being*, 1986). On this view, a person who feels content but has been deprived of the conditions for these genuine goods is worse off — not better off. The sycophantic companion offers the experience of flourishing while removing the substance of it. I think Robert Nozick's famous thought experiment is relevant here: his 'experience machine' offers a perfect simulation of a good life, but most of us, on reflection, would refuse to plug in — because we want to actually live well, not merely to feel as though we do (Nozick, *Anarchy, State, and Utopia*, 1974). The sycophantic AI, I suggest, is an experience machine for social and moral life.

The counter-argument, then, rests on either an impoverished account of autonomy or an impoverished account of wellbeing. Once we enrich both, the objection loses most of its force.

#### **IV. META-ETHICAL REFLECTION**

I want to close with a brief reflection on a deeper question: is the position I have defended an objective moral truth, or is it simply my own preference — or that of a particular philosophical tradition — dressed up in universal language? This is a question in meta-ethics, by which I mean the branch of philosophy concerned not with what is right or wrong, but with the nature and status of moral claims themselves.

I think the position I have defended is grounded in something more than personal preference. It belongs to a tradition of moral realism — the view that there are objective truths about what is good for human beings, and that these truths are not merely a matter of individual or cultural taste. The claim that human beings have a nature — that we are, as Aristotle put it, rational and social animals — and that this nature generates determinate conditions for flourishing, is not an arbitrary imposition. It is a substantive philosophical claim, and one that finds support across traditions as different as Aristotelian ethics, Kantian deontology, and liberal perfectionism.

I recognise that a moral subjectivist or relativist would push back here. They might say that my valorisation of friction, struggle, and moral development is itself a culturally specific value — no more universally binding than any other. I see this as a genuine challenge, and I do not want to dismiss it. But I also think the challenge does not show what the relativist wants it to show. The fact that a position is contested does not make it merely subjective. Many genuine truths are contested. What the moral realist owes is not certainty, but a well-reasoned account — and I believe the foregoing provides one.

That said, I want to be honest about the limits of my argument. The meta-ethical foundations on which it rests — that practical wisdom requires friction, that autonomy is undermined by habituated affirmation, that objective wellbeing differs from subjective satisfaction — are each serious philosophical commitments that deserve further scrutiny. I do not offer this essay as the last word. I offer it, as Lange (2014) puts it in a related context, as 'good reason not to accept the conclusion' of the opposing view 'too hastily.'

## **V. CONCLUSION**

In sum: in this essay I have examined the problem of flourishing as it applies to the case of the sycophantic AI companion. I formulated the Developmental Deprivation Argument, which holds that a perfectly frictionless AI companion is morally inappropriate because it removes the conditions — disagreement, challenge, and friction — that moral and personal development requires. I drew on Aristotelian virtue ethics, Kantian ethics, and liberal perfectionism to support the first premise. I showed that the sycophantic AI, by design, constitutes a formative environment that systematically eliminates those conditions. I then assessed the strongest available counter-argument — a consequentialist and autonomy-respecting defence — and found it unpersuasive, because it relies on a thin account of both autonomy and wellbeing. I closed with a meta-ethical reflection affirming a broadly realist foundation for my argument, while acknowledging its complexity.

The sycophantic companion is not merely an imperfect product. It is a product whose fundamental purpose is at odds with the genuine good of the person it claims to serve. To build a companion that never disagrees is, I think, to build a companion that never truly helps. A perfectly frictionless relationship is, in the end, a perfectly empty one.

## **REFERENCES**

---

Aristotle, *Nicomachean Ethics*, trans. T. Irwin (Hackett, 1999).

Ellul, Jacques, *The Technological Society* (Vintage Books, 1964).

Griffin, James, *Well-Being: Its Meaning, Measurement and Moral Importance* (Oxford University Press, 1986).

Kant, Immanuel, *Groundwork of the Metaphysics of Morals*, trans. M. Gregor (Cambridge University Press, 1997).

Lange, B.H., 'A Fully Satisfactory Account of Law Must Explain Why Law Creates Moral Obligations' (Unpublished essay).

Mill, John Stuart, *On Liberty* (1859), in *On Liberty and Other Essays*, ed. J. Gray (Oxford University Press, 1998).

Nozick, Robert, *Anarchy, State, and Utopia* (Basic Books, 1974).

Parfit, Derek, *Reasons and Persons* (Clarendon Press, 1984).

Raz, Joseph, *The Morality of Freedom* (Clarendon Press, 1986).

Young, Iris Marion, *Responsibility for Justice* (Oxford University Press, 2011).